

Q. Reconfigure the existing deployment front-end and add a port specification named http , exposing port 80/tcp of the existing container nginx.

Q. Create a New service named front-end-svc exposing the container port http

Q. Configure new service to also expose the individual pods via a NodePort on the nodes on which they are scheduled.

=====

Q. Create a pod with nginx image, named as web in dev-test namespace
SOLU

====

```
$ kubectl create ns dev-test
$ kubectl run web --image=nginx --restart=Never -n dev-test
```

=====

Q. show the pods that are part of namespace XYZ , show services that is part of namespace ABC

SOLU

====

```
$ kubectl get po -n XYZ
$ kubectl get svc -n ABC
```

=====

```
kube3a k top po -n kube-system | head -3
NAME                      CPU(cores)   MEMORY(bytes)
coredns-86c58d9df4-562mf    1m          7Mi
coredns-86c58d9df4-zl4mj    1m          7Mi
```

List the top 1 pod with the highest CPU usage

```
kubectl top po -n kube-system | sort -nr -k2 | head -1
                                         ^^<== -k is column number , so
```

column 2 is the CPU

List the top 1 pod with the highest MEM usage

```
kubectl top po -n kube-system | sort -nr -k3 | head -1
                                         ^^<== -k is column , Column 3 is
```

the Memory

1. Create a static pod on a node and put the pod specification file in /etc/kubernetes/manifests.

SOLU

=====

Refer to <https://kubernetes.io/docs/tasks/administer-cluster/static-pod/> , and use the static-web.yaml

```
root@kube3b:/etc/kubernetes/manifests# pwd
/etc/kubernetes/manifests
root@kube3b:/etc/kubernetes/manifests# cat static-web.yaml
apiVersion: v1
kind: Pod
metadata:
  name: static-web
  labels:
    role: myrole
spec:
  containers:
  - name: web
    image: nginx
    ports:
    - name: web
      containerPort: 80
      protocol: TCP
root@kube3b:/etc/kubernetes/manifests#
```

Refer to : <https://medium.com/@sonasingh46/static-pod-in-kubernetes-e3854507655f>

```
# vi /etc/systemd/system/kubelet.service
```

```
[Service]
```

```
ExecStart=/home/kubernetes/bin/kubelet $KUBELET_OPTS \
--pod-manifest-path=/etc/kubernetes/manifests
```

```
# systemctl daemon-reload ; kubectl get po ; if no pod then
# systemctl restart kubelet
```

2. Create a deployment nginx-dns and expose it as service nginx-dns.

Create busybox pod , Use nslookup to get dns result for service & pod and store the values into a file.

SOLU

=====

```
kube3a k run nginx-dns --image=nginx --port=80
```

```

deployment.apps/nginx-dns created
kube3a k expose deployment nginx-dns --port=80 --target-port=80
    service/nginx-dns exposed
kube3a

by default nginx assings label as run=deploymet|pod_name>

kube3a k get deploy,po,svc -l run=nginx-dns -o wide
NAME                               READY   UP-TO-DATE   AVAILABLE   AGE
CONTAINERS   IMAGES   SELECTOR
deployment.extensions/nginx-dns   1/1     1           1           4m44s   nginx-
dns      nginx   run=nginx-dns

NAME                           READY   STATUS    RESTARTS   AGE   IP
NODE             NOMINATED NODE   READINESS GATES
pod/nginx-dns-6b7f858b9b-hzcvc  1/1     Running   0          4m44s   10.244.2.153
kube3c.home.local   <none>           <none>

^^^^^^^^^<<== we need to resolve this for pod.dns

NAME           TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
SELECTOR
service/nginx-dns ClusterIP  10.98.27.246  <none>        80/TCP       3m47s
run=nginx-dns

kube3a
service.dns
      ^^^^^^<<== we need to resolve this for

kube3a k run bbox --image=busybox:1.28 --restart=Never -- /bin/sh -c 'sleep 360000'
pod/bbox created

Resolve pod IP :
kube3a k exec bbox -- nslookup 10.244.2.153 > /tmp/aa
kube3a
Edit /tmp/aa and delete everything except :
  10-244-2-153.nginx-dns.default.svc.cluster.local

kube3a sudo cp /tmp/aa /opt/TTT/pod.dns
kube3a cat /opt/TTT/pod.dns
  10-244-2-153.nginx-dns.default.svc.cluster.local

Next to resolve service for pod :
kube3a k exec bbox -- nslookup nginx-dns > /tmp/bb
kube3a

```

```
Edit /tmp/bb and ensure you ONLY have :  
kube3a cat /tmp/bb  
    nginx-dns.default.svc.cluster.local  
kube3a
```

```
kube3a sudo cp /tmp/bb /opt/TTT/service.dns  
kube3a cat /opt/TTT/service.dns  
    nginx-dns.default.svc.cluster.local  
kube3a
```

=====

3. Create a init container which shouldn't start if a file is not found.

SOLU

====

I created :

```
$ k run bbox --image=busybox --restart=Never --dry-run -o yaml -- /bin/sh -c 'if [  
! -f /opt/XX/xx]  
>then  
>exit 0  
>else  
>sleep 360000000  
>fi' > bbox12.yaml
```

Once I had above then I added entries for volume: and volumeMounts: for emptyDir, then

```
dharmin@kube3a:~$ cat bbox12.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
  creationTimestamp: null  
  labels:  
    run: bbox  
  name: bbox  
spec:  
  volumes:          <<== U need emptyDir: or hostpath:  
  - name: cache-volume <<==  
    emptyDir: {}      <<==  
  containers:  
  - args:  
    - /bin/sh  
    - -c  
    - |-
```

```

if [ ! -f /opt/XX/xx ]
then
exit 0
else
sleep 36000000
fi
image: busybox
name: bbox
volumeMounts:
- mountPath: /opt/XX <== the initcontainer and the actual container shares
same dir
    name: cache-volume
initContainers:
- args:
  - /bin/sh
  - -c
  - 'touch /opt/XX/xx'
image: busybox
name: bbox1
volumeMounts:
- mountPath: /opt/XX <== the initcontainer and the actual container shares
same dir
    name: cache-volume

#
## dnsPolicy: ClusterFirst
## restartPolicy: Never

```

```
kube3a k create -f bbox12.yaml
pod/bbox created
```

```
kube3a k get po --watch
NAME      READY  STATUS      RESTARTS   AGE
bbox      0/1    Init:0/1    0          12s
bbox      0/1    PodInitializing 0          16s
bbox      1/1    Running    0          18s
bbox      0/1    Completed   0          19s
bbox      1/1    Running    1          21s
bbox      0/1    Completed   1          22s
^
```

```
kube3a k logs bbox
FILE PRESENT BOSS
```

```
=====
```

4. Bootstrap a kubernetes node.

<https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet-tls-bootstrapping/>

```
=====
```

5. Use node selector to assign a pod to particular node.

SOLU

====

Label a Node :

```
$ kubectl label no kube3c    node=kube3c
```

Create alpine pod , and assign to a Node labeled node=kube3c :

```
kube3a k run alpine --image=alpine --restart=Never --dry-run -o yaml --command
sleep 36000000 > alpine-nodeselector.yaml
```

```
kube3a cat alpine-nodeselector.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: alpine
    name: alpine
spec:
  containers:
  - command:
    - sleep
    - "36000000"
    image: alpine
    name: alpine
    resources: {}
dnsPolicy: ClusterFirst
restartPolicy: Never
nodeSelector:      <<== edit file and add this manually
  node: kube3c      <<== add this and add this manually
kube3a
```

```
kube3a k create -f  alpine-nodeselector.yaml
pod/alpine created
```

```
kube3a k get po alpine -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED NODE		READINESS GATES				

```
alpine  0/1    ContainerCreating  0          13s    <none>   kube3c.home.local
<none>
<none>
kube3a
```

6. Label a node and it shouldn't be allowed to execute/schedule any pods.

SOLU

====

At present :

```
Taints:           node-role.kubernetes.io/master:NoSchedule
Taints:           <none>
Taints:           <none>
kube3a
```

```
kube3a k taint node kube3c aa=bb:NoSchedule
kube3a k describe no | egrep -i taint
Taints:           node-role.kubernetes.io/master:NoSchedule
Taints:           <none>
Taints:           aa=bb:NoSchedule
kube3a
```

```
kube3a k get po --all-namespaces -o wide | egrep kube3c | wc -l
2
```

Next deploy 20 nginx pods and see if any gets deployed on kube3c :

```
kube3a k get po --all-namespaces -o wide | egrep -i kube3c | wc -l
2
```

AS per above we still have 2 pods that are running on kube3c , as those are part of daemonset

7. Create a deployment name nginx-app with 3 replicas use nginx:1.9.1, do a rolling update to nginx:1.13.5 and then undo it.

New image was asked to use for rolling update.

SOLU

====

```
kube3a k run nginx-app --image=nginx:1.9.1 --replicas=3
deployment.apps/nginx-app created
```

```
kube3a k rollout status deploy nginx-app
```

```
deployment "nginx-app" successfully rolled out
```

```
kube3a k describe deploy nginx-app | egrep -i image:  
  Image:      nginx:1.9.1
```

```
kube3a
```

```
Next update the nginx to ver 1.13.5
```

```
kube3a k set image deployment nginx-app    nginx-app=nginx:1.13.5  
deployment.extensions/nginx-app image updated
```

```
kube3a k rollout status deploy nginx-app
```

```
Waiting for deployment "nginx-app" rollout to finish: 1 out of 3 new replicas have  
been updated...
```

```
Waiting for deployment "nginx-app" rollout to finish: 1 out of 3 new replicas have  
been updated...
```

```
Waiting for deployment "nginx-app" rollout to finish: 1 out of 3 new replicas have  
been updated...
```

```
Waiting for deployment "nginx-app" rollout to finish: 2 out of 3 new replicas have  
been updated...
```

```
Waiting for deployment "nginx-app" rollout to finish: 2 out of 3 new replicas have  
been updated...
```

```
Waiting for deployment "nginx-app" rollout to finish: 2 out of 3 new replicas have  
been updated...
```

```
Waiting for deployment "nginx-app" rollout to finish: 1 old replicas are pending  
termination...
```

```
Waiting for deployment "nginx-app" rollout to finish: 1 old replicas are pending  
termination...
```

```
deployment "nginx-app" successfully rolled out
```

```
kube3a
```

```
kube3a k describe deployment nginx-app | egrep -i image:
```

```
  Image:      nginx:1.13.5
```

```
kube3a
```

```
Undo the upgraqde to version 1 ( previous ) :
```

```
kube3a k rollout undo deployment nginx-app --to-revision=1  
deployment.extensions/nginx-app rolled back
```

```
kube3a k describe deploy nginx-app | egrep -i image:
```

```
  Image:      nginx:1.9.1
```

```
kube3a
```

```
=====
```

8. Node is in not ready status. Find the issue and fix it.

SOLU

=====

ssh onto the worker Node , and run :

```
$ sudo systemctl status -l kubelet
```

If the kubelete service is not runing then :

```
$ sudo systemctl enable kubelet
$ sudo systemctl restart kubelet
```

Finally run below to ensure kubelet service is runing fine :

```
$ sudo systemctl status -l kubelet
```

=====

9. Find all pods which are using same service in a namespace and store the value into a file.

SOLU

=====

Create a ns called boss :

```
kube3a k create ns boss
namespace/boss created
```

```
kube3a k get ns boss
```

NAME	STATUS	AGE
boss	Active	10s

```
kube3a
```

deploy busybox pods and expose port 80 and create service in ns boss.

```
kube3a k get po -n boss
```

NAME	READY	STATUS	RESTARTS	AGE
b	1/1	Running	0	2m27s
b1	1/1	Running	0	2m22s
b2	1/1	Running	0	2m15s
b3	1/1	Running	0	109s
b4	1/1	Running	0	103s
b5	1/1	Running	0	98s
b6	1/1	Running	0	86s

```
kube3a
```

above is total 8 , but then we ignore the header "NAME READY STATUS RESTARTS AGE" ,
so its 7

```
$ sudo echo 7 > /opt/file/X
```

```

kube3a k get svc -n boss
NAME    TYPE      CLUSTER-IP        EXTERNAL-IP      PORT(S)      AGE
b3      ClusterIP  10.96.172.82    <none>          80/TCP       8m
b4      ClusterIP  10.107.213.233  <none>          80/TCP       7m54s
b5      ClusterIP  10.98.80.82    <none>          80/TCP       7m49s
b6      ClusterIP  10.102.82.0    <none>          80/TCP       7m37s
kube3a k get svc -n boss | wc -l
5

```

By ignoring headers , we have 4 :

```
$ sudo echo 4 > /opt/file/XX
```

10. Find all pods with label cpu=loader and sort them by consuption and store the value into a file.

SOLU

====

I have create couple of pods , all have label cpu=loader, lets see which has High CPU Usage
and which consumes more Memory.,,

```

kube3a k top po -l cpu=loader
NAME    CPU(cores)   MEMORY(bytes)
a1      0m           2Mi
a2      1m           4Mi
a3      6m           14Mi
a4      1m           1Mi
a5      0m           0Mi
a6      0m           0Mi
kube3a

```

```

echo "a3" > high.cpu
echo "a3" > high.mem

```

Delete all the above pods that have label cpu: loader

```

kube3a k delete po -l cpu=loader
pod "a1" deleted
pod "a2" deleted
pod "a3" deleted
pod "a4" deleted
pod "a5" deleted
pod "a6" deleted

```

kube3a

=====

11. Find all persistent volumes in a namespace and sort them by capacity and store the value into a file.

SOLU

=====

```
kube3a k get pv --sort-by=.spec.capacity.storage
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS     CLAIM
STORAGECLASS  REASON    AGE
pv-1g        1Gi       RWO          Retain        Available
103s
myvolume     10Gi      RWO, RWX    Retain        Available
4m5s
myvolume1    15Gi      RWO, RWX    Retain        Available
3m11s
pv-1f        100Gi     RWO          Retain        Available
57s
kube3a
```

=====

12. Create a pod with multiple containers including image of redis, nginx , memcached, consul, busybox and alpine.

SOLU

=====

NOTE : ONLY busybox and alpine requires "--command sleep 3600000" , and rest does NOT

```
$ k run pod-all --image=alpine --restart=Never --dry-run -o yaml --command sleep 3600000 > pod-all.yaml
```

```
kube3a cat pod-all.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod-all
  name: pod-all
```

```

spec:
  containers:
    - command:
        - sleep
        - "3600000"
      image: alpine
      name: p1
    #
    - command:
        - sleep
        - "3600000"
      image: busybox
      name: p2
    #
    - image: redis
      name: p3
    #
    - image: memcached
      name: p4
    #
    - image: nginx
      name: p5
    #
    - image: consul
      name: p6
    #
    dnsPolicy: ClusterFirst
    restartPolicy: Never
    #
kube3a
kube3a k create -f pod-all.yaml
pod/pod-all created

After few mins..
kube3a k get po
NAME      READY   STATUS    RESTARTS   AGE
pod-all   6/6     Running   0          4m52s
kube3a
=====
```

13. Create a secret super-secret with credential as bob and create two pods, one which mount secret as volume and another which mount it as env variable.

SOLU

=====

```
kube3a k create secret generic super-secret --from-literal=credential=bob  
secret/super-secret created
```

```
kube3a k get secret super-secret
```

NAME	TYPE	DATA	AGE
super-secret	Opaque	1	76s

```
kube3a k get secret super-secret --export -o yaml
```

```
apiVersion: v1  
data:  
  credential: Ym9i      <<== echo Ym9i | base64 -d , will give you bob  
kind: Secret  
metadata:  
  creationTimestamp: null  
  name: super-secret  
  selfLink: /api/v1/namespaces/default/secrets/super-secret  
type: Opaque
```

kube3a

Create environment variable called TOPSECRET and assign the secret value , refer to
:

<https://kubernetes.io/docs/concepts/configuration/secret/>

```
$ pod-secret.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
  name: secret-env-pod  
spec:  
  containers:  
    - name: mycontainer  
      image: redis  
      env:  
        - name: TOPSECRET  
          valueFrom:  
            secretKeyRef:  
              name: super-secret  
              key: credential
```

In the above secret url , I searched for env: and saw above , which I cp & paste

```
kube3a k create -f pod-secret.yaml  
pod/secret-env-pod created
```

```
kube3a
```

```
kube3a k get po
NAME           READY   STATUS    RESTARTS   AGE
secret-env-pod 1/1     Running   0          80s
kube3a
```

```
kube3a k exec secret-env-pod -- printenv | egrep TOPSECRET
TOPSECRET=bob
kube3a
```

Next we need to copy the secret in a volume :

```
kube3a cat secret-pod.yaml
kind: Pod
apiVersion: v1
metadata:
  name: secret-pod
  labels:
    name: secret-pod
spec:
  volumes:
  - name: secret-volume
    secret:
      secretName: super-secret
  containers:
  - name: secret-pod
    image: nginx
    volumeMounts:
    - name: secret-volume
      readOnly: true
      mountPath: "/etc/secret-volume"
kube3a
```

From the secret url , I searched for volume: , and found above , which I did minor changes and bingo..

```
kube3a k create -f secret-pod.yaml
pod/secret-pod created
kube3a
kube3a k exec secret-pod -- ls -l /etc/secret-volume
lrwxrwxrwx 1 root root 17 Mar 19 22:58 credential -> ..data/credential

kube3a k exec secret-pod -- cat /etc/secret-volume/credential
bob
```

14. Fix a malfunctioning kubernetes cluster.

15. Create a persistent volume with 1GB details provided.

SOLU

====

On k8s.io , seach for "persistentvolume as hostpath" , will see below :

```
$ pv-1G.yaml
kind: PersistentVolume
apiVersion: v1
metadata:
  name: task-pv-volume
  labels:
    type: local
spec:
  storageClassName: manual
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/mnt/data"
```

```
kube3a k create -f pv-1G.yaml
  persistentvolume/task-pv-volume created
```

```
kube3a k get pv
```

NAME	CAPACITY	ACCESS MODES	RECLAIM POLICY	STATUS	CLAIM
STORAGECLASS	REASON	AGE			
task-pv-volume	1Gi	RWO	Retain	Available	
manual		2s			

```
kube3a
```

Q 15.1 - Create a pv of size 1GB in namespace qa , using hostpath /srv/boss

SOLU

====

=====

16. Create a pod nginx and expose it.

SOLU

=====

```
$ k run nginx --image=nginx --restart=Never --port=80 --expose
```

Above will create nginx pod , with port 80 open in firewall , and will expose as ClusterIP

=====

17. Find all nodes in a cluster which have no taints and store the value into a file.

SOLU

=====

```
kube3a k describe no | egrep -i taint
Taints:           node-role.kubernetes.io/master:NoSchedule
Taints:           <none>
Taints:           <none>
kube3a
```

```
$ sudo echo "2" > /opt/XX/k34ft.txt
```

=====

18. Create a namespace website-backend, run jenkins pod inside it and expose it.

SOLU

=====

```
$ k create ns website-backend
$ k run jenkins --image=jenkins --restart=Never -n website-jenkins --port=80
--expose
```

=====

19. ETCD Backup

<https://kubernetes.io/docs/tasks/administer-cluster/configure-upgrade-etcd/>

SOLU

=====

```
kube3a pwd
/home/dharmin/etcd-v3.3.12-linux-amd64
```

```
kube3a ./etcdctl -v
etcdctl version: 3.3.12
API version: 2
kube3a

By default a etcd is in kube-system namespace
kube3a k get po --all-namespaces | awk '$2 ~ /etcd/ {print $2}'
    etcd-kube3a
kube3a

kube3a k get po etcd-kube3a -n kube-system --export -o yaml | egrep "\-file|data-dir|endpont"
- --cert-file=/etc/kubernetes/pki/etcd/server.crt <<== NEED THIS
- --data-dir=/var/lib/etcd <<== NEED THIS
- --key-file=/etc/kubernetes/pki/etcd/server.key <<== NEED THIS
- --peer-cert-file=/etc/kubernetes/pki/etcd/peer.crt
- --peer-key-file=/etc/kubernetes/pki/etcd/peer.key
- --peer-trusted-ca-file=/etc/kubernetes/pki/etcd/ca.crt
- --trusted-ca-file=/etc/kubernetes/pki/etcd/ca.crt <<== NEED THIS
kube3a

kube3a ./etcdctl backup --help
NAME:
  etcdctl backup - backup an etcd directory

USAGE:
  etcdctl backup [command options]

OPTIONS:
  --data-dir value      Path to the etcd data dir      <<== NEED THIS
  --wal-dir value       Path to the etcd wal dir
  --backup-dir value    Path to the backup dir      <<== NEED THIS
  --backup-wal-dir value Path to the backup wal dir
  --with-v3             Backup v3 backend data

kube3a

kube3a sudo ./etcdctl --cert-file /etc/kubernetes/pki/etcd/server.crt --key-file
/etc/kubernetes/pki/etcd/server.key \
    --ca-file /etc/kubernetes/pki/etcd/ca.crt backup --data-dir
/var/lib/etcd --backup-dir /tmp/ABC
[sudo] password for dharmin:
2019-03-20 00:05:36.754230 I | ignoring v3 raft entry
2019-03-20 00:05:36.754479 I | ignoring v3 raft entry
```

```
2019-03-20 00:05:36.754856 I | ignoring v3 raft entry
2019-03-20 00:05:36.755113 I | ignoring v3 raft entry
```

Lets check out backed-up data :

```
kube3a sudo chmod -R 755 /tmp/ABC
kube3a
```

```
kube3a find /tmp/ABC -print
/tmp/ABC
/tmp/ABC/member
/tmp/ABC/member/wal
/tmp/ABC/member/wal/0000000000000000-0000000000000000.wal
/tmp/ABC/member/snap
/tmp/ABC/member/snap/00000000000007-0000000000a3973.snap
/tmp/ABC/member/snap/db
kube3a
```

Lets check if the backed up data is good :

On k8s.io , search for ETCDCTL_API=v3 , and in the pg you see scroll down and shold see :

```
kube3a ETCDCTL_API=3 ./etcdctl --write-out=table snapshot status
/tmp/ABC/member/snap/db
+-----+-----+-----+-----+
| HASH    | REVISION | TOTAL KEYS | TOTAL SIZE |
+-----+-----+-----+-----+
| 7fb25e0b |        0 |          1 |      25 kB |
+-----+-----+-----+-----+
kube3a
```

20. Create nginx as daemonset.

SOLU

=====

Remember a DaemonSet is similar to a Deployment, so create a Deployment manifest and change is to a daemonset :

```
kube3a k run redis-ds --image=redis --restart=Always --dry-run -o yaml > redis-ds.yaml
^~~~~~<<== this is optional, don't
```

```
require
```

```
$ vi redis-ds.yaml
apiVersion: apps/v1
kind: DaemonSet      <<== this needs to be changed from Deployment to DaemonSet
metadata:
  creationTimestamp: null
  labels:
    run: redis-ds
  name: redis-ds
spec:
  replicas: 1  <<===== Delete this
  selector:
    matchLabels:
      run: redis-ds
  strategy: {}  <<===== Delete this
  template:
    metadata:
      creationTimestamp: null
      labels:
        run: redis-ds
    spec:
      containers:
        - image: redis
          name: redis-ds
          resources: {}  <<== delete this
status: {}  <<= Delete this
```

```
kube3a k create -f  redis-ds.yaml
daemonset.apps/redis-ds created
```

```
kube3a k get ds,po  -o wide
NAME                           DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE
NODE SELECTOR     AGE   CONTAINERS   IMAGES   SELECTOR
daemonset.extensions/redis-ds  2        2         2        2           2
<none>            35s   redis-ds     redis     run=redis-ds
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED NODE	READINESS	GATES				
pod/redis-ds-45dd5	1/1	Running	0	35s	10.244.2.152	
kube3c.home.local	<none>		<none>			
pod/redis-ds-tf5qp	1/1	Running	0	35s	10.244.1.164	
kube3b.home.local	<none>		<none>			

```
kube3a
```

=====

21. Scale number of pods of a deployment to 6.

SOLU

=====

```
$ k run nginx --image=nginx --replicas=2  
$ k scale deployment nginx --replicas=6
```

=====

22. Store the logs of a container having error file-not-found.

SOLU

=====

```
$ k log <pod_name> | egrep "file-not-found" > log1  
$ sudo cp log1 /opt/XXX/kdrvsv
```

=====

23. Create a pod with local volume details provided.

SOLU

=====

refer to k8s.io and search for hostPath or emptyDir

=====

24. Create a pod and assign it to node labelled as disk=spinning.

SOLU

=====

check if the nodes have the label disk=spinning :

```
$ k get no ---show-labels
```

\$ k run redis --image=redis --restart=Never --dry-run -o yaml > redis-pod.yaml
then at the end of the line add :

```
  nodeSelector:  
    disk: spinning
```

Ensure the "nodeSelector:" is in line with "containers:"

=====

```
if use busybox or alpine pod then always do use --command sleep 3600000 or --  
/bin/sh -c 'command'
```

when comes to nginx or redis pod's , does not require --command or -- /bin/sh ,
as per above

PRE-REQ

=====

```
$ vi a
alias k=kubectl
alias c=clear

$ source ~/a
$ cp a /tmp
```

When you become root , then :

```
# source /tmp/a
```

0. Create following deployments with 2 pods : nginx, redis, memcached, consul
expose nginx to service boss
Forget above question

Can you tell which pods are part of service boss ?

SOLU

==

If you create a service from a deployment/pod , the default label for service is same as of deployment/pod.

Below will show you Selector , which is same as label :

```
$ kubectl get svc boss --export -o yaml
```

Once you know what label is assinged to above service , then

```
$ kubectl get po -l <label_name>
```

You'll see the pods that use the service boss.

=====

1. Create nginx static pod , copy the yaml file in /etc/kubernetes/manifest dir.

SOLU

==

On worker Node , as user root :

```
# source /tmp/a
# cd /etc/kubernetes/manifests
copy the static pod here from https://kubernetes.io/docs/tasks/administer-cluster/static-pod/
```

```
# systemctl status -l kubelet
kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset:
disabled)
   Drop-In: /etc/systemd/system/kubelet.service.d
             10-kubeadm.conf
     Active: active (running) since Mon 2019-03-11 11:47:55 GMT; 6s ago
       Docs: https://kubernetes.io/docs/home/
Main PID: 1624 (kubelet)
```

Our worker Node kube3c is online (Ready) :

```
kube3a k get no
```

NAME	STATUS	ROLES	AGE	VERSION
kube3a	Ready	master	5d2h	v1.13.0
kube3b	Ready	<none>	5d1h	v1.13.0
kube3c	Ready	<none>	5d1h	v1.13.0
kube3a				

14. If you need to know what options to use with securityContext in a pod , how can you find out

SOLU

====

```
$ kubectl explain pod.spec.securityContext --recursive
```

15. Assign a Worker Node with taint tz=tanzania:NoSchedule and assign nginx pod to a node

with the taint tz=tanzania:NoSchedule.

SOLU

====

Show current taint :

```
$ k describe no | egrep -i taints:
Taints:           node-role.kubernetes.io/master:NoSchedule
Taints:           <none>
Taints:           <none>
```

Add a taint to worker node kube3c :

```
$ k taint node kube3c tz=tanzania:NoSchedule
      node/kube3c tainted
```

```

$ k describe no | egrep -i taints:
Taints:           node-role.kubernetes.io/master:NoSchedule
Taints:           <none>
Taints:           tz=tanzania:NoSchedule
$

k run nginx --image=nginx --restart=Never --dry-run -o yaml > nginx-pod1.yaml
$
$ cat nginx-pod1.yaml
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
  name: nginx
spec:
  containers:
  - image: nginx
    name: nginx
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Never
  tolerations:             <<== added this
  - key: "tz"               <<== added this
    operator: "Equal"       <<== added this
    value: "tanzania"       <<== added this
    effect: "NoSchedule"   <<== added this
$
$ k apply -f nginx-pod1.yaml
pod/nginx created

$ k get po -o wide
NAME      READY   STATUS    RESTARTS   AGE     IP          NODE   NOMINATED NODE
READINESS GATES
nginx    1/1     Running   0          43s    192.168.80.221   kube3c   <none>
<none>
$
```

AS seen above the pod is deployed on kube3c , which has the taint.

Do the clean up Now..

```

$ k describe no kube3c | egrep -i taint
  Taints:           tz=tanzania:NoSchedule
$ k taint node kube3c tz-
```

```

node/kube3c untainted
$ k describe no kube3c | egrep -i taint
    Taints:           <none>
$
Deleted the nginx pod
=====
```

16. Refer to <https://kubernetes.io/docs/tasks/debug-application-cluster/core-metrics-pipeline/>

PRE-REQ : 1. Install Metric Server on the Master Node

2. Run below commands on Master Node :

app=hello	kubectl run nginx --image=nginx	--restart=Never -l
app=hello	kubectl run redis --image=redis	--restart=Never -l
app=hello	kubectl run consul --image=consul	--restart=Never -l
app=hello	kubectl run memcached --image=memcached	--restart=Never -l

Q. From the pods label app: hello , name the pod that shows high CPU and Memory usage ,

dump the output to /tmp/cpu.usage and /tmp/mem.usage

SOLU

```

$ k top po -l app=hello
NAME          CPU(cores)   MEMORY(bytes)
consul        14m          23Mi
memcached    1m           5Mi
nginx         0m           2Mi
redis         2m           9Mi
$
```

```

$ k top po -l app=hello | sort -nr -k2 | head -1 | awk '{print $1}' | sudo tee
/tmp/cpu.usage
                                         ^^<<== refer to 2nd Column
```

```

$ k top po -l app=hello | sort -nr -k3 | head -1 | awk '{print $1}' | sudo tee
/tmp/mem.usage
```

^^<<== refer to 3rd Column

I am not sure if they want ONLY the pod name or.....

You can :

```
$ k top po consul | tail -1 | awk '{print $1}' > /tmp/cpu.usage , then edit and  
only leave 1 pod entry with higher cpu usage  
$ k top po consul | tail -1 | awk '{print $1}' > /tmp/cpu.memory , then edit and  
only leave 1 pod entry with higher Mem usage
```

=====

17. Sort the PV in ascending order from top to bottom.

SOLU

====

```
$ k get pv --sort-by=.spec.capacity.storage  
NAME      CAPACITY      ACCESS MODES      RECLAIM POLICY      STATUS      CLAIM  
STORAGECLASS   REASON      AGE  
all-1      2609192632320m    RWX          Retain          Available    manual  
65s  
all-2      2780991324160m    RWX          Retain          Available    manual  
65s  
all-3      2791728742400m    RWX          Retain          Available    manual  
65s  
all-4      3006477107200m    RWX          Retain          Available    manual  
65s  
all-6      6Gi            RWX          Retain          Available    manual  
65s  
all-5      6227702579200m    RWX          Retain          Available    manual  
65s  
$  
$
```

17.a Sort all below pv by name

```
NAME      CAPACITY      ACCESS MODES      RECLAIM POLICY      STATUS      CLAIM  
STORAGECLASS   REASON      AGE  
myvolume      10Gi        RWO, RWX      Retain          Available  
normal        6s          RWO, RWX      Retain          Available  
myvolume1     15Gi        RWO, RWX      Retain          Available  
normal        5s          RWO, RWX      Retain          Available  
pv-1f         100Gi       RWO          Retain          Available  
5s  
pv001-1      2Gi         RWO, RWX      Retain          Available
```

```

manual           6s
  task-pv-volume   1Gi        RWO, RWX      Retain       Available
manual           6s

SOLU
====

kube3a k get pv --sort-by=.metadata.name
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS  CLAIM
STORAGECLASS  REASON    AGE
myvolume      10Gi     RWO, RWX      Retain       Available
normal
myvolume1     15Gi     RWO, RWX      Retain       Available
normal
pv001-1       2Gi      RWO, RWX      Retain       Available
manual
pv-1f         100Gi    RWO          Retain       Available
95s
task-pv-volume 1Gi      RWO, RWX      Retain       Available
manual
96s
kube3a
=====
```

```

18. From etcd pod  get the log error msg "store.index" and copy to /opt/KK/xx
SOLU
====

$ k log etcd-kube3a -n kube-system | egrep "store.index" | sudo tee
/opt/KK/xx
=====
```

```

19. Create nginx deployment that has 2 pods,  with servicename as nginx-svc ,
exposed thru nodeport.
  Resolve Pod IP and Service Name in DNS in any container.
SOLU
====
```

```

$ k run nginx --image=nginx --replicas=2 --port=80
$ k expose deployment nginx --type=NodePort --port=80 --target-port=80 --
name=nginx-svc
=====
```

After I executed above :

```
$ k get deploy,svc
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE	
deployment.extensions/nginx	2/2	2	2	53s	
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	6d5h
service/nginx-svc	NodePort	10.103.112.224	<none>	80:30071/TCP	19s

We know the default label for nginx pod is run=nginx , so lets use the default label and

The best way to get pod IP is :

```
$ k get po -l run=nginx -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED NODE	READINESS	GATES				
nginx-57867cc648-7n2tc	1/1	Running	0	11m	192.168.79.226	kube3b
<none>	<none>					
nginx-57867cc648-bm2vf	1/1	Running	0	11m	192.168.80.222	kube3c
<none>	<none>					

IP's ^^^^^^^^^^<<== Pod

Resolve the pod IP's in consul pod :

```
$ k exec consul -- nslookup 192.168.79.226
```

Name: 192.168.79.226

Address 1: 192.168.79.226 192-168-79-226.nginx-svc.default.svc.cluster.local

\$

```
$ k exec consul -- nslookup 192.168.80.222
```

Name: 192.168.80.222

Address 1: 192.168.80.222 192-168-80-222.nginx-svc.default.svc.cluster.local

\$

Resolve the Service Name , the service name is shown when run k get svc :

```
$ k exec consul -- nslookup nginx-svc
```

Name: nginx-svc

Address 1: 10.103.112.224 nginx-svc.default.svc.cluster.local

\$

REF : <https://medium.com/@jmarhee/using-initcontainers-to-pre-populate-volume-data-in-kubernetes-99f628cd4519>

20. Include busybox initcontainer that creates empty host file in /data , use non persistent volume.

```
apiVersion: v1
```

```
kind: Pod
```

```

metadata:
  name: my-app
spec:
  containers:
  - name: my-app
    image: nginx
    volumeMounts:
    - name: config-data
      mountPath: /data
#
# SOLU
=====
I created in steps , in first step I made sure the above part is able to create a
pod :
$ cat nginx-pod2.yaml
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  volumes:
  - name: config-data
    emptyDir: {}
  containers:
  - name: my-app
    image: nginx
    volumeMounts:
    - name: config-data
      mountPath: /data
$
```

Above pod worked fine..

Next was to create another busybox pod that would create empty hosts file in /data , where /data is emptyDir

```

$ k run bbox --image=busybox --restart=Never --dry-run -o yaml -- /bin/sh -c
'touch /data/hosts' > bbox1.yaml
I extracted info from bbox1.yaml and create my final pod manifest.
$ cat nginx-pod2.yaml
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  volumes:
```

```

- name: config-data
  emptyDir: {}
containers:
- name: my-app
  image: nginx
  volumeMounts:
  - name: config-data <== this matches the name under volume:
    mountPath: /data
# below is from busybox container
initContainers:
- args:
  - /bin/sh
  - -c
  - touch /data/hosts
image: busybox
name: bbox
volumeMounts:
- name: config-data <== this matches the name under volume:
  mountPath: /data
# Above is from busybox container
dnsPolicy: ClusterFirst
restartPolicy: Never

```

After I created above pod :

```

$ k get po
NAME      READY     STATUS    RESTARTS   AGE
my-app    1/1       Running   0          10m

```

```

$ k exec my-app -- ls /data/hosts
/datas/hosts

```

```

$ k exec my-app -- ls -l /data/hosts
-rw-r--r-- 1 root root 0 Mar 13 14:30 /data/hosts
$ 

```

Time to delete my pod :

```

$ k delete po my-app
pod "my-app" deleted
$ 
=====
```

REF : <https://raw.githubusercontent.com/openshift-evangelists/kbe/master/specs/ic/deploy.yaml>

21 . Create alpine init container so can be used with deploy , do create path to

```
local disk.  
apiVersion: apps/v1beta1  
kind: Deployment  
metadata:  
  name: ic-deploy  
spec:  
  replicas: 1  
  template:  
    metadata:  
      labels:  
        app: ic  
    spec:  
      containers:  
      - name: main  
        image: centos:7  
        command:  
        - "bin/bash"  
        - "-c"  
        - "while true; do cat /ic/this; sleep 5; done"  
      volumeMounts:  
      - mountPath: /ic  
        name: msg
```

SOLU

=====

Loking at above manifest, seems like volume is missing , so lets add that , and run : k run alpine --image=alpine --restart=Never --dry-run -o yaml -- /bin/sh -c 'echo BOSSSS > /ic/this' > alpine.yaml
use info from alpine.yaml file and add in below deployment.

```
$ cat ic-deploy.yaml  
apiVersion: apps/v1beta1  
kind: Deployment  
metadata:  
  name: ic-deploy  
spec:  
  replicas: 1  
  template:  
    metadata:  
      labels:  
        app: ic  
    spec:  
      volumes:  
      - name: msg  
        hostPath:
```

```

    path: /mnt/data
  containers:
  - name: main
    image: centos:7
    command:
    - "bin/bash"
    - "-c"
    - "while true; do cat /ic/this; sleep 5; done"
  volumeMounts:
  - mountPath: /ic
    name: msg
  initContainers:
  - args:
    - /bin/sh
    - -c
    - echo BOSSS > /ic/this
    image: alpine
    name: alpine
    volumeMounts:
    - mountPath: /ic
      name: msg
##      dnsPolicy: ClusterFirst    <== this is part of the alpine pod,
##      restartPolicy: Never       <== was generated when created alpine pod
$  

$ k get deploy,po
NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.extensions/ic-deploy      1/1     1            1           32m
  

NAME                                READY   STATUS        RESTARTS   AGE
pod/ic-deploy-786f5d879d-pckxg    1/1     Running       0          32m
$  

$ k exec ic-deploy-786f5d879d-pckxg -- cat /ic/this
    BOSSS
$  

$ k logs ic-deploy-786f5d879d-pckxg
    BOSSS
  

=====

```

22. Add initContainer and non persistent disk

```

apiVersion: v1
kind: Pod
metadata:

```

```
name: happypanda
spec:
  containers:
    - name: busybox
      image: busybox:latest
      command: ["/bin/sh", "-c"]
      args: ["cat /opt/workdir/helloworld && sleep 3600"]
      volumeMounts:
        - name: workdir
          mountPath: /opt/workdir
```

SOLU

=====

```
$ k run alpine --image=alpine --restart=Never --dry-run -o yaml -- /bin/sh -c
'echo BOSSS > /opt/workdir/helloworld'

$ cat initcon.yaml
apiVersion: v1
kind: Pod
metadata:
  name: happypanda
spec:
  volumes:
    - name: workdir
      emptyDir: {}    <<== this is non persistent disk
#
  containers:
    - name: busybox
      image: busybox
      command: ["/bin/sh", "-c"]
      args: ["cat /opt/workdir/helloworld && sleep 3600"]
      volumeMounts:
        - name: workdir
          mountPath: /opt/workdir
#
# Below was generated from above "k run alpine ...."
  initContainers:
    - args:
        - /bin/sh
        - -c
        - echo BOSSS > /opt/workdir/helloworld
      image: alpine
      name: alpine
      volumeMounts:           <<== This I copied from above
```

```

- name: workdir           <<== "XXXXXXXXXXXXXXXXXXXX"
  mountPath: /opt/workdir <<== "XXXXXXXXXXXXXXXXXXXX"
##
## dnsPolicy: ClusterFirst <<== this is part of the alpine pod , don't require
## restartPolicy: Never    <<== "XXXXXXXXXXXXXXXXXXXX"
$
```

After I deployed above pod , then :

```
$ k get po happypanda
```

NAME	READY	STATUS	RESTARTS	AGE
happypanda	1/1	Running	0	4m17s

I ran "k describe po happypanda" , and only showing you few lines..

Events:

Type	Reason	Age	From	Message
Normal	Scheduled	5m14s	default-scheduler	Successfully assigned default/happypanda to kube3c
Normal	Pulling	5m13s	kubelet, kube3c	pulling image "alpine"
Normal	Pulled	5m6s	kubelet, kube3c	Successfully pulled image "alpine"
Normal	Created	5m6s	kubelet, kube3c	Created container
Normal	Started	5m6s	kubelet, kube3c	Started container
Normal	Pulling	5m6s	kubelet, kube3c	pulling image "busybox"
Normal	Pulled	5m4s	kubelet, kube3c	Successfully pulled image "busybox"
Normal	Created	5m4s	kubelet, kube3c	Created container
Normal	Started	5m4s	kubelet, kube3c	Started container

```
$
```

23. List a Node with high CPU Usage and copy the output to /opt/TTT/cpu.higg.usage
List a Node with High Memory Usage and save the output to
/opt/TTT/mem.usage.high

SOLU

BECOME ROOT , then execute the kubectl config user-config <cluster_name> , then :
if kubectl command does not work then may need to :
KUBECONFIG=/etc/kubernetes/admin.conf

```
root@kube3a:~# kubectl top no
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
kube3a	73m	3%	1167Mi	14%
kube3b.home.local	15m	0%	326Mi	4%
kube3c.home.local	16m	0%	355Mi	4%

```

root@kube3a:~#
root@kube3a:~# kubectl top no | sort -rn -k2 | head -1 | awk '{print $1}'
      kube3a
                           ^^^<<== -k is the column no.

root@kube3a:~# kubectl top no | sort -rn -k2 | head -1 | awk '{print $1}' >
/opt/TTT/cpu.higg.usage
root@kube3a:~# cat /opt/TTT/cpu.higg.usage
      kube3a
root@kube3a:~#

root@kube3a:~# kubectl top no | sort -rn -k4 | head -1 | awk '{print $1}'
      kube3a
root@kube3a:~# kubectl top no | sort -rn -k4 | head -1 | awk '{print $1}' >
/opt/TTT/mem.usage.high

root@kube3a:~# cat /opt/TTT/mem.usage.high
      kube3a

```

24. Count number of Nodes that are in ready state , except ones which shows tolerance : noschedule

save the results in /opt/TTT/node-ready

SOLU

====

As we don't have perm to wrie in /opt/TTT dir , for that we'll need to become

root :

```

$ cp a /tmp/a
kube3a sudo -i
root@kube3a:~#
root@kube3a:~# source /tmp/a
kube3a
kube3a id -a
      uid=0(root) gid=0(root) groups=0(root)
kube3a

```

copy & paste from top of the Q : kubectl config use-config <cluster_name>

kube3a k get no

The connection to the server localhost:8080 was refused - did you specify the right host or port?

kube3a export KUBECONFIG=/etc/kubernetes/admin.conf

```
kube3a
```

```
kube3a k get no
```

NAME	STATUS	ROLES	AGE	VERSION
kube3a	Ready	master	14d	v1.13.1
kube3b.home.local	Ready	<none>	14d	v1.13.1
kube3c.home.local	Ready	<none>	9d	v1.13.1

```
kube3a k describe no | egrep -i taint:
```

```
kube3a k describe no | egrep -i taint
```

```
Taints: node-role.kubernetes.io/master:NoSchedule
```

```
Taints: <none>
```

```
Taints: <none>
```

```
kube3a
```

```
kube3a echo 2 > /opt/TTT/node-ready
```

```
kube3a cat /opt/TTT/node-ready
```

```
2
```

```
kube3a
```

```
25. You are unable to Query the cluster nodes :
```

```
root@kube3a:~# kubectl get no
```

```
The connection to the server localhost:8080 was refused - did you specify  
the right host or port?
```

```
SOLU
```

```
=====
```

```
If kubelet service is down you can also get above error. Run :
```

```
systemctl status -k kubelet
```

```
You need to check 2 files , as the /etc/{BROKEN|broken} could be in either of the  
file.
```

```
$ sudo vi /etc/systemd/system/kubelet.service
```

```
OR
```

```
$ sudo vi /lib/systemd/system/kubelet.service
```

```
REPLACE : /etc/[BROKEN|broken] OR /etc/kuernetes/BROEKN
```

```
WITH : /etc/kubernetes/manifests
```

```
$ sudo systemctl daemon-reload
```

```
$ systemctl status -l kubelet
```

```
OR
```

```
$ systemctl restart kubelet
```

```
Also check that the kubelete service is enabled :  
$ systemctl list-unit-files | egrep -i kubelet
```

```
Here you'll need to wait for say 2mins , before able to run command kubectl get no
```

```
=====
```

```
26. Create a deployment which contains one nginx container. The deployment needs a  
label
```

```
called app:frontend and should be configured for 5 replicas.
```

```
Expose a deployment with clusterIP service that listens on tcp/8080 and target  
tcp/8080.
```

```
SOLU
```

```
====
```

```
=====
```

```
27.
```

```
1. create a pod named "web" using image nginx:1.11.9-alpine, on ports 80 and 443  
SOLU
```

```
=====
```

```
kube3a k run web --image=nginx:1.11.9-alpine --restart=Never --port=80 --dry-run  
-o yaml > nG.yaml
```

```
kube3a
```

```
kube3a vi nG.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  creationTimestamp: null  
  labels:  
    run: web  
  name: web  
spec:  
  containers:  
  - image: nginx:1.11.9-alpine  
    name: web  
    ports:  
    - containerPort: 80  
    - containerPort: 443      <<== I manually added this entry  
    resources: {}  
dnsPolicy: ClusterFirst
```

```
restartPolicy: Never
status: {}
```

```
kube3a k create -f nG.yaml
pod/web created
```

```
kube3a k get po
NAME READY STATUS RESTARTS AGE
web 1/1 Running 0 15s
kube3a
```

```
kube3a k describe po web | egrep -i port
Ports: 80/TCP, 443/TCP
Host Ports: 0/TCP, 0/TCP
kube3a
```

2. create a service to expose that pod, named as "webservice"
SOLU
====

```
kube3a k expose pod web --port=80,443 --name=webservice
service/webservice exposed
```

```
kube3a k get svc webservice
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
webservice ClusterIP 10.107.250.6 <none> 80/TCP, 443/TCP 8s
kube3a
```

3. copy the dns records for the service in file /opt/TTT/web.dnsrecord
record the pod dns in /opt/TTT/pod.dns

SOLU

====

```
kube3a k run bbox --image=busybox:1.28 --restart=Never -- /bin/sh -c 'sleep 36000'
pod/bbox created
```

```
kube3a k get po -o wide
NAME READY STATUS RESTARTS AGE IP NODE
NOMINATED-NODE READINESS GATES
bbox 1/1 Running 0 45s 10.244.1.10 kube3b.home.local <none>
<none>
web 1/1 Running 0 11m 10.244.2.25 kube3c.home.local <none>
<none>
kube3a
```

^^^^^^^^^<== this is the pod IP , which

```
needs to resolve as pod dns
```

```
kube3a k exec bbox -- nslookup 10.244.2.25 | sudo tee /opt/TTT/pod.dns
[sudo] password for dharmin:
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local
```

```
Name: 10.244.2.25
Address 1: 10.244.2.25 10-244-2-25.webservice.default.svc.cluster.local
kube3a
```

```
kube3a k get svc
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1      <none>        443/TCP      21d
webservice  ClusterIP  10.107.250.6    <none>        80/TCP,443/TCP  12m
kube3a
```

Here we resolving DNS name of service webservice :

```
kube3a k exec bbox -- nslookup webservice | sudo tee /opt/TTT/web.dnsrecord
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local
```

```
Name: webservice
Address 1: 10.107.250.6 webservice.default.svc.cluster.local
kube3a
```

One i found which wasn't there was

Finding all nodes which are in ready status (except those where tolerance : noschedule)

Also - they answer that Sanjay told to restore the faulty node also worked

Sol - open kubelet systemd file n you will see /etc/ BROKEN will be written

Just remove broken n add /Kubernetes/manifests

Do daemon reload n kubelet restart

Note - this will take time
My advice to u will be do write ASAP
All questions are same
I'm sure you will pass
One more thing - we have to only find out ONE node that has HIGH cpu utilization
Not all the nodes

=====

Kubectl get node OR pod - you'll see Connection Refused Error.

One Master and One Worker Node.

Login to the FIRST Node

Open the Kubelet Service File.

sudo systemctl status kubelet

Open this file : /lib/systemd/system/kubelet.service

Similar to the

-pod-manifest-path

In that file, you'll see

/etc/broken

You've to change is this to /etc/kubernetes/manifest

=====

NOTE :

For metric server you have to include below :

- command:

- /metrics-server
- --kubelet-insecure-tls
- --kubelet-preferred-address-types=InternalIP
- --logtostderr